

Lab 2 Assignment

```
library(ISLR)
library(MASS)
library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.1 --

## v ggplot2 3.3.5      v purrr 0.3.4
## v tibble 3.1.6       v dplyr 1.0.7
## v tidyr 1.1.4        v stringr 1.4.0
## v readr 2.1.1        v forcats 0.5.1

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## x dplyr::select() masks MASS::select()

library(caret)

## Loading required package: lattice

##
## Attaching package: 'caret'

## The following object is masked from 'package:purrr':
##
## lift
```

Lab Assignment

Q1. Discuss this example with your group. (no submission needed)

Galton's data on the heights of parents and their children, by child.

This data set lists the individual observations for 934 children in 205 families on which Galton (1886) based his cross-tabulation.

Suppose you are tasked with building a machine learning algorithm that predicts the *son's height* Y, using the *father's height* X.

```
library(HistData)

##GaltonFamilies
head(GaltonFamilies)
```

##	family	father	mother	midparentHeight	children	childNum	gender	childHeight
## 1	001	78.5	67.0	75.43	4	1	male	73.2
## 2	001	78.5	67.0	75.43	4	2	female	69.2
## 3	001	78.5	67.0	75.43	4	3	female	69.0
## 4	001	78.5	67.0	75.43	4	4	female	69.0
## 5	002	75.5	66.5	73.66	4	1	male	73.5
## 6	002	75.5	66.5	73.66	4	2	male	72.5

```
table(GaltonFamilies$family)
```

```
##
## 001 002 003 004 005 006 007 008 009 010 011 012 013 014 015 016
## 4 4 2 5 6 1 6 3 1 1 8 1 2 2 3 9
## 017 018 019 020 021 022 023 024 025 026 027 028 029 030 031 032
## 6 3 1 8 3 3 7 1 2 5 3 6 3 1 6 5
## 033 034 035 036 037 038 039 040 041 042 043 044 045 046 047 048
## 5 1 5 4 4 6 2 5 1 6 2 2 3 8 4 3
## 049 050 051 052 053 054 055 056 057 058 059 060 061 062 063 064
## 7 2 2 5 9 4 5 5 5 7 1 2 4 6 1 5
## 065 066 067 068 069 070 071 072 073 074 075 076 077 078 079 080
## 1 11 4 5 8 5 6 7 3 2 7 7 4 5 8 1
## 081 082 083 084 085 086 087 088 089 090 091 092 093 094 095 096
## 4 9 8 4 5 4 4 4 8 7 3 2 4 2 3 5
## 097 098 099 100 101 102 103 104 105 106 107 108 109 110 111 112
## 10 1 8 3 4 6 7 4 6 7 9 7 7 4 1 3
## 113 114 115 116 117 118 119 120 121 122 123 124 125 126 127 128
## 1 6 7 3 1 3 5 11 8 4 5 9 3 4 1 2
## 129 130 131 132 133 134 135 136 136A 137 138 139 140 141 142 143
## 3 11 2 2 7 4 8 10 8 4 5 1 10 8 4 1
## 144 145 146 147 148 149 150 151 152 153 154 155 156 157 158 159
## 4 8 6 1 1 5 1 2 1 5 1 7 4 1 10 5
## 160 161 162 163 164 165 166 167 168 169 170 171 172 173 174 175
## 1 8 6 5 4 3 11 4 8 3 5 1 8 9 5 6
## 176 177 178 179 180 181 182 183 184 185 186 187 188 189 190 191
## 8 5 1 2 6 7 1 4 1 15 4 1 4 5 9 2
## 192 193 194 195 196 197 198 199 200 201 202 203 204
## 6 6 2 3 4 5 7 7 1 2 2 3 2
```

```
set.seed(1983)
galton_heights <- GaltonFamilies %>%
  filter(gender == "male") %>%
  group_by(family) %>% #grouped data frame
  sample_n(1) %>%
  ungroup() %>% #removes grouping
  select(father, childHeight) %>%
  rename(son = childHeight)
```

```
head(galton_heights)
```

```
## # A tibble: 6 x 2
##   father son
##   <dbl> <dbl>
## 1  78.5  73.2
## 2  75.5  72.5
## 3   75   71
## 4   75  70.5
## 5   75   69
## 6   74  76.5
```

(i) Let's split the data into training and testing.

```
#?slice
#slice() lets you index rows by their (integer) locations.
```

#It allows you to select, remove, and duplicate rows.

```
y <- galton_heights$son
test_index <- createDataPartition(y, times = 1, p = 0.3, list = FALSE) #30% for the test set

train_set <- galton_heights %>% slice(-test_index)
test_set <- galton_heights %>% slice(test_index)
```

(ii) The *best guess(prediction)* would be the mean;

```
(yhat <- mean(train_set$son))
```

```
## [1] 69.14839
```

then the error would be (yhat-ybar)

```
mean((yhat - test_set$son)^2)
```

```
## [1] 7.738353
```

(iii) Can we do better? Yes.

$$f(x) = E(Y|X = x) = \beta_0 + \beta_1 x$$

```
fit <- lm(son ~ father, data = train_set)
fit$coef
```

```
## (Intercept)      father
##  37.6101619    0.4570757
```

$$\hat{f}(x) = 43.7 + 0.36x$$

We can see that this does indeed provide an improvement over our guessing approach.

```
y_hat <- fit$coef[1] + fit$coef[2]*test_set$father
mean((y_hat - test_set$son)^2)
```

```
## [1] 6.634095
```

(iv) Predict

```
y_hat <- predict(fit, test_set)
mean((y_hat - test_set$son)^2)
```

```
## [1] 6.634095
```

Learn more;

```
?predict.lm
?predict.glm
```

many machine learning algorithms have a predict function.

Q2: Coefficient Interpretation

Suppose we have the following model to estimate GPA for Georgetown undergraduate students.

$$GPA = \beta_0 + \beta_1 * Female + \beta_2 Age$$

Where *Age* is a continuous variable measured in years and gender is encoded as:

- Female = 1 if student is female
 - Female = 0 if student is male
- a) What is the interpretation of the β_0 coefficient?
 - b) What is the interpretation of the β_1 coefficient?
 - c) Suppose $\beta_0 = 1.5$ and $\beta_1 = .5$ and $\beta_2 = 1.5$. What is the expected GPA of a student who is male and 20 years old?
 - d) Suppose you believe that males have lower GPAs than females. Further, suppose you believe that GPA tends to decrease as students get older. What would this imply about the signs of the coefficients of the model?
 - e) Suppose we change the units of measurement for *Age* from years to days. What effect would that have on β_2 ? This is an important reminder about interpreting the absolute magnitude of the coefficients.

Question 3

Data Science Question: Are the variables contributing for predicting “Valence” of the songs is same for both musicians; Taylor Swift and John Legend?

Valence: A measure from 0.0 to 1.0 describing the musical positiveness conveyed by a track. Tracks with high valence sound more positive (e.g. happy, cheerful, euphoric), while tracks with low valence sound more negative (e.g. sad, depressed, angry).

```
Artists <- read.csv("Artists.csv",header = TRUE)
head(Artists)
```

```
##   X  artist_name Valence danceability energy loudness speechiness acoustictness
## 1 1 Taylor Swift  0.328      0.594  0.713   -5.314    0.0503    0.000328
## 2 2 Taylor Swift  0.408      0.516  0.777   -4.908    0.0375    0.001080
## 3 3 Taylor Swift  0.299      0.645  0.593   -6.506    0.0288    0.034400
## 4 4 Taylor Swift  0.767      0.584  0.557   -6.371    0.0342    0.012900
## 5 5 Taylor Swift  0.132      0.440  0.528   -7.809    0.0317    0.017100
## 6 6 Taylor Swift  0.642      0.642  0.695   -5.620    0.0281    0.000443
##   liveness  tempo                                track_name
## 1   0.1140 129.958                State Of Grace (Taylor's Version)
## 2   0.0761 125.047                                Red (Taylor's Version)
## 3   0.1300 109.984                Treacherous (Taylor's Version)
## 4   0.0576 154.008 I Knew You Were Trouble (Taylor's Version)
## 5   0.2340 185.972                All Too Well (Taylor's Version)
## 6   0.0753 103.984                                22 (Taylor's Version)
##   album_name album_release_year
## 1 Red (Taylor's Version)        2021
## 2 Red (Taylor's Version)        2021
## 3 Red (Taylor's Version)        2021
## 4 Red (Taylor's Version)        2021
## 5 Red (Taylor's Version)        2021
## 6 Red (Taylor's Version)        2021
```

Create a new variable with the name "danceable" and add it to the “Artists” data set. Fill this variable with values “High”, “Medium”, “Low” according to the danceability value such that 0.8-1 is “High”, 0.5-0.79 is “Medium” and 0.0-0.49 is “Low”.

```
Artists$danceable<-rep(0, nrow(Artists))
```

```

Artists1 <- within(Artists, {

  danceable[danceability>= 0.800 & danceability <= 1] <- "High"
  danceable[danceability>= 0.500 & danceability < 0.8] <- "Medium"
  danceable[danceability <= 0.499] <- "Low"
} )

head(Artists1)

##   X  artist_name Valence danceability energy loudness speechiness acoustiness
## 1 1 Taylor Swift  0.328      0.594  0.713   -5.314    0.0503    0.000328
## 2 2 Taylor Swift  0.408      0.516  0.777   -4.908    0.0375    0.001080
## 3 3 Taylor Swift  0.299      0.645  0.593   -6.506    0.0288    0.034400
## 4 4 Taylor Swift  0.767      0.584  0.557   -6.371    0.0342    0.012900
## 5 5 Taylor Swift  0.132      0.440  0.528   -7.809    0.0317    0.017100
## 6 6 Taylor Swift  0.642      0.642  0.695   -5.620    0.0281    0.000443
##   liveness  tempo                                track_name
## 1   0.1140 129.958                State Of Grace (Taylor's Version)
## 2   0.0761 125.047                                Red (Taylor's Version)
## 3   0.1300 109.984                Treacherous (Taylor's Version)
## 4   0.0576 154.008 I Knew You Were Trouble (Taylor's Version)
## 5   0.2340 185.972                All Too Well (Taylor's Version)
## 6   0.0753 103.984                                22 (Taylor's Version)
##   album_name album_release_year danceable
## 1 Red (Taylor's Version)          2021   Medium
## 2 Red (Taylor's Version)          2021   Medium
## 3 Red (Taylor's Version)          2021   Medium
## 4 Red (Taylor's Version)          2021   Medium
## 5 Red (Taylor's Version)          2021    Low
## 6 Red (Taylor's Version)          2021   Medium

```

Just do the analysis for John Legend.

- Fit a linear regression models separately for Artists: *John Legend* , to predict *Valence* of the songs on all the quantitative variables in the data set and “danceable” variable you created earlier .

Remember to train your model on a training set and validate the model using a validation set.

- Check the multicollinearity c. In cooperate interaction terms c. what is the best model? with or without the ineractions?